

# Using the GNU assembler for Intel processors

## 1 Introduction

This tutorial is a minimal description of what you need to know to generate code for an 80386 (or later) Intel processor, using the *as* [3] assembler on a Linux system. The appendix contains instructions and modifications for running the example programs using the C library, instead of the Linux kernel, for input and output.

There are two kinds of assemblers:

1. Intel style assemblers. The main assembler of this kind is *nasm*.
2. AT&T style assemblers, mainly the Gnu assembler, *as* or *gas*.

The main differences between them are the order of the operands in the instructions, the special characters indicating constants, registers, contents of memory locations, operand types ([3], Section 8.8). They both generate the same machine instructions, so the choice is mostly a matter of taste.

Most modern processors have 64-bit registers for addressing and arithmetics. Most computers in the computer rooms of the E-building have such registers, but all of them are running in 32-bit mode. If you are using a computer running in 64-bit mode, you have to modify your code according to Section 13.

There is a lot of documents on the web describing how to program this processor. A good one, using the Intel style, is by Carter [6].

## 2 The processor

There are eight 32-bit general purpose registers: **eax**, **ebx**, **ecx**, **edx**, **ebp**, **edi**, **esi**, and **esp**. The last one, **esp**, should be reserved for use as a stack pointer. The other registers may be used freely, but some instructions require the use of specific registers. The letters **e** and **x** in the first four registers just indicates that they are 32 bits long. Conventionally, the base pointer register, **ebp**, is used as a pointer to the current activation record, and **edi** and **esi** are used as destination and source pointers by string instructions.

The first four registers may also be used as byte and 16-bit word registers. The name of the last 8 bits of **eax** is **al**, and the last 16 bits may be referred to by **ax**. The first 8 bits of **ax** are referred to by **ah**. There are analogous names for parts of the **ebx**, **ecx** and **edx** registers. Additionally, there are one bit flags that may be set by some instructions and used by other instructions.

The processor can execute more than 300 different instructions. This document describes about 30 of them. Each instruction has at most two operands. There are four kinds of operands: registers, memory addresses, constant data held by the instruction, or implicit values contained in registers. With few exceptions, an instruction can refer to at most one memory location.

### 3 An example

In the below example, the number of digits needed to print a non-negative number is computed by repeatedly dividing the number by 10 until it becomes 0.

```

        .data                                # allocating memory
n:       .long    234                        # the number
length:  .long    0                          # the result
ten:     .long    10                         # the divisor

        .text                                # instructions
        .global _start                      # make _start globally known
_start:  movl     $0, %ebx                    # use ebx as counter
        movl     n, %eax                     # copy number to eax
nextdigit:
        movl     $0, %edx                    # prepare for long division
        idivl    ten                         # divide combined edx:eax registers by 10
                                                # quotient to eax, remainder to edx

        addl     $1, %ebx                     # add 1 to counter
        cmpl     $0, %eax                     # compare eax to 0
        jg       nextdigit                   # jump if eax>0
        movl     %ebx, length                 # copy counter to memory
                                                # exit to OS kernel to terminate execution
        movl     $0, %ebx                     # first argument: exit code
        movl     $1, %eax                     # sys_exit index
        int      $0x80                       # kernel interrupt
```

The program has two *sections*, a `.data` section that describes how to allocate memory for global variables, and a `.text` section with the instructions. The sections may appear in any order. Each instruction should start a new line. No indentation is required, but improves readability. Comments start with a `#` character.

Each variable has a label, a type, and an initial value. The label is a name that denotes the memory address of the variable. All variables in this example have the type `.long` which uses 32 bits. The variable `n` is the number to analyse (234), the number of digits will be stored in the variable `length` (3), and the variable `ten` is the divisor.

The first line of the `.text` section is a *directive* making a label accessible outside this section. `_start` is the default label used by the loader for the first instruction to be executed.

All register references start with a percent character. The first instruction,

```
movl $0, %ebx
```

sets the `ebx` register to 0. `$0` is a *constant* operand. The dollar sign is important, because without it the value at location 0 in memory will be used. In the example, we use the `ebx` register to count the digits.

The next two `movl` instructions prepare for a division of 234 by 10. The first one

```
movl n, %eax
```

copies the value, at the memory location with label `n`, to the `eax` register.

Division is performed on the combined 64 bits of the `edx` and `eax` registers. We set the first register to 0. The `idivl` instruction performs signed integer division, i.e., it assumes that negative numbers are represented as two complements. The `idivl` instruction has one operand, the divisor. This instruction cannot take constant operands, so the value of the operand is fetched from memory location `ten`. After the execution of the `idivl` instruction, the quotient ends up in `eax` and the remainder in `edx`. We add 1 to `ebx` to count this digit. The actual value of the digit is in `edx`.

Next, 0 is compared to `eax`. Notice the order of the operands! If the contents of `eax` is greater than 0, the next instruction (`jg nextdigit`) makes a jump to `nextdigit`. After another two divisions, the comparison prevents the jump from happening. Finally, the value in `ebx` is saved in the memory location `length`.

The proper way to terminate the execution is to call the `exit` procedure in the operating system kernel using an *interrupt*. The last three lines do that with an exit code equal to 0 signaling normal return.

The program may now be translated to machine code by the assembler, *as*, linked by *ld*, and executed with no visible result.

```
> as -o digit1.o digit1.s
> ld -o digit1 digit1.o
> ./digit1
>
```

## 4 Debugging with ddd

We may use the debugger *ddd* to inspect registers and memory during the execution. *ddd* is a graphical user interface to the Gnu debugger, *gdb*. You should invest some time in learning to use it.

The assembler requires an extra option to keep the symbol table in the executable program.

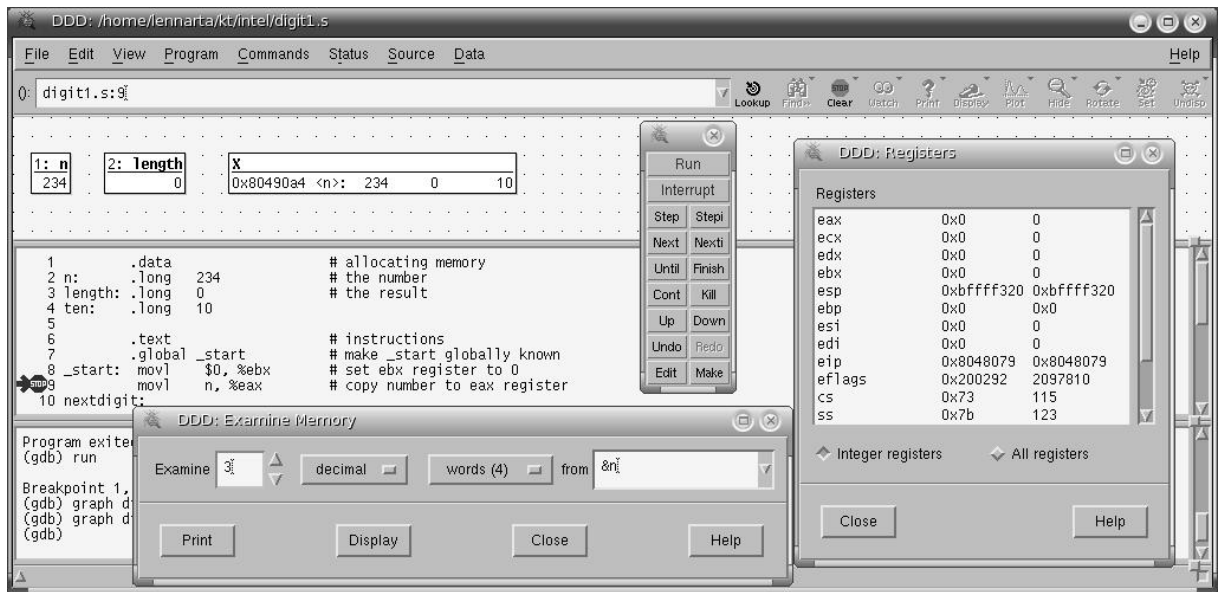
```
> as --gstabs -o digit1.o digit1.s
> ld -o digit1 digit1.o
> ddd digit1&
```

In *ddd*, breakpoints are inserted and deleted by right clicking on the line number and selecting the appropriate item. A breakpoint at the first instruction will have no effect.

The buttons **Run**, **Step**, and **Cont** are used to execute the program, single step the execution, and continue the execution after a break point.

Variables are displayed by right clicking on the label.

The Register window is opened via **Status**→**Registers**, and sections of the memory are displayed through the **Data**→**Memory** dialog.



## 5 Memory allocation

In programming languages supporting recursion, all variables will be stored in activation records on a stack. When Linux is used on an Intel-based machine, there is no need to allocate a stack because the linker, *ld*, will reserve 2 Mb of memory for this purpose. The stack pointer register, *esp*, will point to the top of this stack.

Constant and global data may be allocated in the `.data` section. Some examples:

```
.data
v32:    .long 0           # 32 bits, initial value 0
v16:    .word 0xffff      # 16 bits, all ones
v8:     .byte '-'         # 8 bits, ascii code for -
vs:     .ascii "input"    # string with 5 bytes
vs0:    .asciz "input"    # string with 6 bytes, last byte is 0
        .align 4         # align at 32-bit word boundary
stack:  .skip 1024        # allocate 1024 bytes
tos = .                   # tos will be the address of the next byte
```

Most C functions expect strings to be terminated by a 0 byte generated by `.asciz`.

The directive `.align 4` allocates the next item at a 4\*8 bit address boundary. Memory accesses to 32-bit words should have this alignment.

At the label `stack`, 1024 bytes are allocated.

## 6 Operands

An instruction can have four kinds of operands:

- **Constant:** A constant is preceded by a dollar sign, for example, `$10` or `$n`. The last operand denotes the address corresponding to the label `n`.

- **Register:** A register operand starts with a % character.
- **Address:** An address operand refers to a value at a memory location. The most common example is a label of a variable or an instruction.
- **Implicit:** Some instructions have implicit operands, e.g. the `idivl` instruction uses the `eax` and `edx` registers.

Examples of how values and instructions can be referenced, where `length`, `nextdigit`, and `stack` refer to names from previous examples:

Operand	Refers to
<code>length</code>	value at label <code>length</code>
<code>length+4</code>	value at 4 bytes after <code>length</code>
<code>length-4</code>	value at 4 bytes before <code>length</code>
<code>nextdigit</code>	instruction at <code>nextdigit</code>
<code>(%ebp)</code>	value at address contained in <code>ebp</code>
<code>4(%ebp)</code>	value at 4 bytes after address contained in <code>ebp</code>
<code>stack(%eax)</code>	value at <code>stack+eax</code>
<code>(%ebp,%eax,4)</code>	value at <code>ebp+4*eax</code>
<code>stack(%ebp,%eax,4)</code>	value at <code>stack+ebp+4*eax</code>

Note, that the assembler can evaluate simple arithmetic expressions.

## 7 Instructions

The following table lists the most common instructions for arithmetic operations and copying of data. For each instruction, the permitted kinds of operands are indicated by the initial letters of register, memory, and constant, together with the number of bits taking part in the operation. There are similar instructions for byte (8 bits) and word (16 bits) operands. The trailing 1 in instruction names is then replaced by `b` or `w`.

Instruction	Operands	Effect
<code>movl</code>	<code>rmc32, rm32</code>	<code>rm32 = rmc32</code>
<code>addl</code>	<code>rmc32, rm32</code>	<code>rm32 = rm32+rmc32</code>
<code>subl</code>	<code>rmc32, rm32</code>	<code>rm32 = rm32-rmc32</code>
<code>negl</code>	<code>rm32</code>	<code>rm32 = -rm32</code>
<code>incl</code>	<code>rm32</code>	<code>rm32 = rm32+1</code>
<code>decl</code>	<code>rm32</code>	<code>rm32 = rm32-1</code>
<code>imull</code>	<code>rmc32, r32</code>	<code>r32 = r32*rmc32</code>
<code>imull</code>	<code>rm32</code>	<code>edx:eax = r32*eax, 64 bit result</code>
<code>idivl</code>	<code>rm32</code>	<code>eax = edx:eax/rm32, edx = remainder</code>
<code>notl</code>	<code>rm32</code>	<code>rm32 = ! rm32, bitwise, false = 0</code>
<code>andl</code>	<code>rmc32, rm32</code>	<code>rm32 = rm32 &amp; rmc32, bitwise</code>
<code>orl</code>	<code>rmc32, rm32</code>	<code>rm32 = rm32   rmc32, bitwise</code>
<code>cmpl</code>	<code>rmc32<sub>1</sub>, rmc32<sub>2</sub></code>	compare by computing <code>rmc32<sub>2</sub>-rmc32<sub>1</sub></code>
<code>leal</code>	<code>m32, r32</code>	<code>r32 = location denoted by m32</code>

Note, that an instruction can have at most one address operand.

The `leal` instruction computes the address of an operand and saves the result in a register.

The `cmpl` instruction (and some other arithmetic instructions) set flags that can be used by conditional instructions. For instance, a jump may be unconditional or conditional:

Instruction	Effect
<code>jmp dest</code>	jump unconditionally
<code>jg dest</code>	jump if greater than
<code>jle dest</code>	jump if less than

The above is not a complete list of all jump instructions. That is, the right side of the jump instruction name, after the `j`, may be replaced by either of the following *conditional codes*:

Conditional codes					
<b>l</b>	<b>le</b>	<b>e</b>	<b>ne</b>	<b>g</b>	<b>ge</b>
<	≤	=	≠	>	≥


Creating additional jump instructions, like `jle` or `je`.

These conditional codes (*cc*) may also be used by other instructions. For instance, depending on whether a condition holds or not, a byte can be set to 1 or 0, or a word can be copied:

Instruction	Effect
<code>setcc rm8</code>	$rm8 = cc ? 1 : 0$
<code>cmovcc rm32, r32</code>	$r32 = rm32$ if <i>cc</i>

## 8 Stack instructions

As mentioned above, the Linux linker (*ld*) allocates a stack intended for activation record, and sets the stack pointer register (`esp`) to the top of the stack. The stack grows towards the bottom of the memory, i.e. to the left in the figure below. The address in the stack pointer is indicated by a thick line. (`%esp`) refers to the topmost value on the stack.



Pushing and popping of 32-bit operands is done with the following instructions:

Instruction	Operand	Effect
<code>pushl</code>	<code>rmc32</code>	push value in <code>rmc32</code>
<code>popl</code>	<code>rm32</code>	pop to <code>rm32</code>

When using the calling convention used by the C compiler, the `ebx` register should be restored to its original value after a procedure call. Hence, this value should be pushed on the stack on entry to a procedure:

```
pushl %ebx
```

### Stack

					ebx value	value
--	--	--	--	--	-----------	-------

Likewise, at the end of the procedure the value should be restored. Assuming that the value of the stack pointer is the same, **ebx** is restored by:

`popl %ebx`

### Stack

					ebx value	value
--	--	--	--	--	-----------	-------

When a program is executed in an operating system, the execution may be *interrupted* by events in the operating system. During such an interrupt, other machine instructions, which use registers and the stack, may be executed. Before these instructions are allowed to run, the operating system performs a *context switch*, where the state of the current program is saved. When the operating system transfers control back to the program, the saved state is restored and execution continues.

## 9 Procedure calls

There are some instructions that support procedure calls. Most compilers respect the conventions used by the C compiler, which makes it easier to call procedures compiled by different compilers.

Instruction	Operands	Effect
<b>call</b>	dest	push return address and jump
<b>ret</b>		pop return address and jump
<b>ret</b>	c32	pop return address and c32 bytes
<b>int</b>	c32	interrupt to kernel

The instruction pointer, **eip**, holds the address of the next instruction to be executed. The jump instructions modify this register, but it cannot be modified directly by a **movl** or an arithmetic instruction.

When a **call dest** instruction is executed, the address of the next instruction is pushed onto the stack, and a jump to the instruction at the **dest** label is performed.

To return from a procedure, the **ret** instructions should be used. This instruction assumes that the return address is at the top of the stack, pops it and jumps to that location.

The **ret c32** instruction pops the return address and the specified number of bytes from the stack, and proceed execution at the address. This instruction may be used to deallocate the memory used by the arguments of the calling procedure.

The C compiler expects the arguments to be pushed onto the stack, in the activation record of the calling procedure, before pushing the return address. The arguments should be pushed in reverse order.

The following diagrams show what happens to the stack during a procedure call with two arguments. Before the call, the stack pointer (**esp**) points to the topmost word currently in use, indicated by a thick line.

### Stack

						current frame
--	--	--	--	--	--	---------------

After pushing of two arguments, the situation is as follows.

```
pushl arg2
pushl arg1
```

### Stack

				arg1	arg2	current frame
--	--	--	--	------	------	---------------

Next we call the procedure, p

```
call p
```

The return address is pushed onto the stack:

### Stack

			ret adr	arg1	arg2	current frame
--	--	--	---------	------	------	---------------

Assuming that the base pointer (**ebp**) is used to indicate the start of the current activation record, it should be saved on the stack in order to be restored upon return from the procedure. The base pointer is reassigned to make it easy to find the arguments and local variables.

```
pushl %ebp
movl %esp, %ebp
```

### Stack

		dyn link	ret adr	arg1	arg2	current frame
--	--	----------	---------	------	------	---------------

The first argument may now be accessed with `8(%ebp)`. At the end the procedure, restore the stack pointer and returns:

```
popl %ebp
ret
```

### Stack

				arg2	arg1	current frame
--	--	--	--	------	------	---------------

After the return the arguments must be popped.

```
addl $8, %esp
```

### Stack

						current frame
--	--	--	--	--	--	---------------

In the following example, the code from page 2 has been extended to print the value of the number and use two procedures. The first procedure computes a string representation of a given non-negative number and prints it using the OS kernel and the second one just returns control to the kernel.



```

        .text
        .global _start
_start:
        pushl    $234                # push argument
        call     writeint
        addl     $4, %esp            # pop stack
        pushl    $0                  # push argument
        call     exit

writeint:
        pushl    %ebp                # save old base pointer
        movl     %esp, %ebp          # set base pointer
        movl     8(%ebp), %eax        # copy argument to eax
        movl     $10, %ebx           # set divisor to 10
        subl     $12, %esp           # allocate for result string
        movl     %ebp, %edi          # edi points to previous digit
writedigit:
        movl     $0, %edx            # divide edx:eax ...
        idivl    %ebx                # by 10
        addl     $'0', %edx          # convert remainder to ascii
        decl     %edi                # push ...
        movb     %dl, (%edi)         # digit
        cmp      $0, %eax            # jump if eax>0
        jg       writedigit          # let the OS kernel print the string
        movl     %ebp, %edx          # compute ...
        subl     %edi, %edx          # third argument: string length
        movl     %edi, %ecx          # second argument: string address
        movl     $1, %ebx            # first argument: file descriptor
        movl     $4, %eax            # sys_write call index
        int      $0x80              # kernel interrupt
        addl     $12, %esp           # deallocate result string
        popl     %ebp                # restore base pointer
        ret                          # return
exit:
        movl     4(%esp), %ebx       # first argument: error code
        movl     $1, %eax            # sys_exit call index
        int      $0x80              # kernel interrupt

```

Since these procedures do not use global variables their activation records have no static links.

If you execute this example, it may happen that the printed result will not be visible, since the shell prompt overprints the result. Use `less` or redirection to avoid this.

The file `eda180.s` shown in the appendix includes versions of `printint` and `readint` that can handle negative numbers, and some more procedures that may be useful in the course project. They respect the C conventions, restoring the values of the `ebp`, `ebx`, and `edi` registers upon `exit`.

## 10 Representation of memory words

The Intel architecture uses a peculiar representation for a word in memory called *little endian*. It means that the bytes within a word are stored in reverse order.

You may observe this when inspecting the same part of the memory using hexadecimal bytes and hexadecimal words. Looking at string data using words the character will appear in reverse order within each word. While, if you inspect 32-bit integer data using bytes the least significant byte will appear first.

This might be confusing, but can usually be ignored.

## 11 Static links

There are two instructions that support block structured languages by setting up a *display* of static links that makes access to global variables efficient.

Instruction	Operands	Effect
<b>enter</b>	c32, c5	set up dynamic and c5 static links, allocate c32 bytes
<b>leave</b>		deallocate ditto and restore <b>ebp</b>

The **enter** instruction will push the contents of the **ebp** register and follow the statics links c5 times and push each link on the stack. Then it will allocate c32 bytes for local variables by decrementing **esp**. The second argument should be equal to the static nesting level of the procedure. The main procedure should have level 1.

The **leave** instruction will deallocate the memory used for the links and local variables and restore **ebp**.

The following pseudo code describes in detail the execution of

**enter** c32, c5

when  $c5 \geq 1$  using a temporary variable **frameptr**.

```
pushl %ebp
movl %esp, frameptr
repeat (c5-1) times {
    subl 4, %ebp
    pushl (%ebp);
}
pushl frameptr
movl frameptr, %ebp
subl c32, %esp
```

The **leave** instruction performs

```
movl %ebp, %esp
popl %ebp
```

There is no illustrating example using these instructions since we believe that it is more instructive to handle the static links by yourself, but you are free to explore this alternative. Notice that the offset of the first variable in a frame depends on the frame level.

## 12 Interfacing C functions

It is easy to call C functions from an assembler program. The following program uses three functions from the standard C library. It reads a number from the keyboard using `scanf`, adds 1, prints the result using `printf`, and terminates by calling `exit`.

```
.text
.global main
main:
    pushl    $n                # push second arg, address of n
    pushl    $sfmt             # push first arg, address of sfmt
    call     scanf              # call scanf("%d", &n)
    addl     $8, %esp           # pop 2 arguments
    addl     $1, n
    pushl    n                  # push second argument, n
    pushl    $fmt               # push first argument, address of fmt
    call     printf             # call printf("%d\n", eax)
    addl     $8, %esp           # pop 2 arguments
    pushl    $0                 # push first argument, exit code = 0
    call     exit               # call exit(0)

.data
n:      .long    0              # number
fmt:    .asciz   "%d\n"         # format for printf
sfmt:   .asciz   "%d"           # format for scanf
```

The `scanf` function has two arguments. The first argument, `"%d"`, is the format string describing how the input string should be interpreted as a decimal number. The second argument should be the address where the resulting number should be saved.

The `printf` function requires one or more arguments. The first argument is the memory address of a string describing how to print the other arguments. In this case, we use the string `"%d\n"` to indicate that we shall print a decimal number followed by a newline character. The second argument is the number to be printed.

We terminate the execution by calling the `exit` function with 0 as an argument. Since this call will not return it is useless to restore the stack pointer in a subsequent instruction.

The standard C library is a *shared library*. This means that several processes using the same library function can share one copy in memory. There is some overhead for this that you should leave to the compiler. The compiler can take an assembler program as input. It will generate a small program with the `_start` label and a call to a procedure called `main`. This should be the first label in your assembler program.

By convention, a C function will use the `eax` register when returning a value. You should assume that all registers except `esp`, `ebp`, `ebx`, and `edi` may have changed. You may “compile” the file `writen.s` using the `gcc` (or `cc`) command and execute the program with the following option to use `ddd`:

```
> gcc -gstabs -o writen writen.s
> ./writen
234
235
```

A program compiled by the C compiler may call assembler functions. It will then expect the values in `ebx`, `esi`, `edi`, `esp`, and `ebp` to be unchanged upon return.

You may compile a C program and inspect the generated assembler instructions using the `-S` option: `gcc -S -o main.s main.c`

## 13 64-bit mode

When using 64-bit mode, the names of the general registers have an initial `r` instead of `e`: `rax`, `rbx`, `rcx`, `rdx`, `rbp`, `rdi`, `rsi`, and `rsp`. The 32 rightmost bits of the first four registers are still available using their old names: `eax`, `ebx`, `ecx`, and `edx`. Most instructions have a 64 bit with `l` replaced by `q`. Exceptions are `push` and `pop`. An example program:

```
.text
.global main
.global _start
_start:
    call main
    movl $0, %ebx
    movl $1, %eax
    int $0x80
main:
    pushq %rbp
    movq %rsp, %rbp
    subq $8, %rsp
    movl $1, 16(%rbp)
    movq %rbp, %rsp
    popq %rbp
    ret
```

For all the details, see [2]

## 14 Mac OS

Appendix 14 contains an example that may be assembled using `gcc` under Mac OS.

When calling external functions compiled by `gcc`, like `printf` and `exit`, the stack must be aligned on a 16-byte boundary. This means that the last hexadecimal digit of `%esp` must be 0 when the call instruction is executed. Otherwise you will get a **Segmentation error**. See e.g. [4] for further details.

It may be best to let your compiler generate code to assure proper alignment.

In order to use the program `digit.s`, the directive `.global` must be changed to `.globl` and the label `_start` must be replaced by `start` at both occurrences.

## References

- [1] Intel, Manuals, [developer.intel.com/design/Pentium4/documentation.htm#manuals](http://developer.intel.com/design/Pentium4/documentation.htm#manuals)
- [2] Intel, Manuals, [www.intel.com/products/processor/manuals/](http://www.intel.com/products/processor/manuals/)
- [3] Gnu, *Using as*, [gnu-mirror.dkuug.dk/software/binutils/manual/gas-2.9.1/as.html](http://gnu-mirror.dkuug.dk/software/binutils/manual/gas-2.9.1/as.html)
- [4] Sanglard, F., *IA-32 assembly on Mac OS X*, <http://www.fabiensanglard.net/macosxassembly/index.php>
- [5] Zeller, A., *Debugging with ddd*, [www.gnu.org/manual/ddd/pdf/ddd.pdf](http://www.gnu.org/manual/ddd/pdf/ddd.pdf)
- [6] Carter, P.A., *PC Assembly Language*, [www.drmpaulcarter.com/pcasm/](http://www.drmpaulcarter.com/pcasm/)
- [7] The Netwide Assembler, [sourceforge.net/projects/nasm](http://sourceforge.net/projects/nasm)

# Appendix. eda180.s

Printed by Lennart And

Oct 21, 07 15:19	eda180.s	eda180.s	Page 1/2
<pre> .text .global writeint ## void writeint(int n) ## writes n on sysout  .global readint ## int readint() ## reads an integer from sysin ## return value is in eax  .global writestr ## void writestr(char[] str, int n) ## writes n characters from str on sysout  .global writeln ## void writeln() ## writes a newline character on sysout  .global exit ## void exit(int n) ## terminates execution with error code n  writeint:     pushl    %ebp     pushl    %ebx     pushl    %edi     movl     \$0, %ecx     movl     16(%ebp), %eax     movl     \$10, %ecx     subl     \$12, %esp     movl     %ebp, %edi     cmpl     \$0, %eax     jge      writedigit     negl     %eax  writedigit:     movl     \$0, %edx     idivl    %ecx     addl     \$'0', %edx     decl     %edi     movb     %dl, (%edi)     cmpl     \$0, %eax     jge      writedigit     negl     %eax      # save old base pointer     # save ebx     # save edi     # set base pointer     # copy argument to eax     # set divisor to 10     # allocate for local string     # set edi     # argument negative?     # negate     # set up for 64 bit division     # divide edx:eax by 10     # ascii digit in edx     # push ...     # digit     # quotient=0 ?     # argument negative?     # push ...     # ',-'     # third argument: string length     # second argument: string address     # first argument: file descriptor     # sys.write interrupt index     # kernel interrupt sys_write     # restore stack pointer     # restore edi     # restore ebx     # restore base pointer     # return     # save old base pointer     # save ebx     # save edi  readint:     pushl    %ebp     pushl    %ebx     pushl    %edi </pre>	<pre>     movl     %esp, %ebp     subl     \$12, %esp     movl     %esp, %edi     movl     \$12, %edx     movl     %edi, %ecx     movl     \$0, %ebx     movl     \$3, %eax     int      \$0x80     addl     %eax, %ecx     decl     %ecx     # parse input     movl     \$0, %eax     cmpl     \$'-', -12(%ebp)     jne      readdigit     incl     %edi     readdigit:         imull    \$10, %eax         movl     \$0, %edx         movb     (%edi), %dl         subl     \$'0', %edx         addl     %edx, %eax         incl     %edi         cmpl     %edi, %ecx         jg       readdigit         cmpl     \$'-', -12(%ebp)         jne      rest         negl     %eax         movl     %ebp, %esp         popl     %edi         popl     %ebx         popl     %ebp         ret      rest:         movl     %ebx         movl     8(%esp), %edx         movl     %1, %ebx         movl     \$4, %eax         int      \$0x80         popl     %ebx         ret      writestr:         pushl    %ebx         movl     12(%esp), %edx         movl     8(%esp), %ecx         movl     %1, %ebx         movl     \$4, %eax         int      \$0x80         popl     %ebx         ret      writeln:         pushl    %ebx         pushl    \$10         movl     \$1, %edx         movl     %esp, %ecx         movl     \$1, %ebx         movl     \$4, %eax         int      \$0x80         popl     %edi         popl     %ebx         popl     %ebp         ret      exit:         movl     4(%esp), %ebx         movl     \$1, %eax         int      \$0x80 </pre>	<pre> # set base pointer # allocate memory # maximal string length # string address # file descriptor # sys.read interrupt index # kernel interrupt sys_read # address of ... # last character # result accumulator # negative number? # skip ',-' # multiply result by 10 # copy digit ... # to dl # convert ascii to value # add value to result # address of next digit # end of string? # negative number? # negate result # restore stack pointer # restore edi # restore ebx # restore base pointer # save ebx # third argument: string length # second argument: string address # first argument: file descriptor # sys_write interrupt index # kernel interrupt sys_write # restore ebx # return # save ebx # push newline character # third argument: string length # second argument: string address # first argument: file descriptor # sys_write interrupt index # kernel interrupt sys_write # pop stack # restore ebx # return # first argument: error code # sys_exit interrupt index # kernel interrupt sys_exit </pre>	Page

# Appendix. Using the C library

Printed by Lennart And

Apr 20, 06 9:56	digit1c.s	Page 1/1
<pre>.data      234 .long      0 length: .long      0 ten:      .long     10  .text .global main .extern exit  main:      movl \$0, %ebx            movl n, %eax            # use ebx as counter            # copy number to eax  nextdigit: # prepare for long division            # divide combined edx:eax registers by 10            # quotient to eax, remainder to edx            idivl ten            addl \$1, %ebx            cmpl \$0, %eax            jg nextdigit            # jump if eax&gt;0            movl %ebx, length            # copy counter to memory            # terminate execution            # call C function exit to terminate execution            pushl \$0            call exit            # first argument: exit code = 0            # call exit(0)  # compile and debug using: # gcc -g -o digit1c digit1c.s # ddd digit1c&amp;</pre>	<pre># allocating memory # the number # the result # the divisor  # use ebx as counter # copy number to eax  # prepare for long division # divide combined edx:eax registers by 10 # quotient to eax, remainder to edx # add 1 to counter # compare eax to 0 # jump if eax&gt;0 # copy counter to memory # terminate execution # first argument: exit code = 0 # call exit(0)  # compile and debug using: # gcc -g -o digit1c digit1c.s # ddd digit1c&amp;</pre>	

Apr 20, 06 9:58	digit2c.s	Page
<pre>.text .global main .extern printf .extern exit  main:      pushl \$234            call writeint            addl \$4, %esp            push \$0            call exit  writeint:   %esp, %ebp            movl 4(%ebp), %eax            movl \$10, %ebx            decl %esp            movb \$0, (%esp)  nextdigit: # divide edx:eax ...            movl \$0, %edx            idivl %ebx            addl \$'0, %edx            decl %esp            movb %dl, (%esp)            cmp \$0, %eax            jg nextdigit             # call printf(esp) to print the string            pushl %esp            call printf            addl \$4, %esp            movl %ebp, %esp            ret  # compile and run using: # gcc -o digit2c digit2c.s # ./digit2c&amp;</pre>	<pre># save stack pointer # copy argument to eax # divisor # push ... # null character # divide edx:eax ... # by 10 # convert remainder to ascii # push ... # digit # jump if eax&gt;0  # push address of string # pop argument (unnecessarily) # restore stack pointer</pre>	

# Appendix. Using the C library under Mac OS

Printed by Lennart And

Mar 12, 10 10:04	digitmac.s	Page 1/1
<pre>.data buffer: .asciz " " .text .globl _main _main: main:     nop     pushl \$234     call writeint     addl \$4,%esp     subl \$0x8,%esp     pushl \$0     call _exit  writeint:     movl %esp,%ebp     movl 4(%ebp),%eax     movl \$10,%ebx     movl \$buffer,%edi     addl \$16,%edi     decl %edi     movb \$0, (%edi)  nextdigit:     movl \$0,%edx     idivl %ebx     addl \$'0,%edx     decl %edi     movb %dl, (%edi)     cmp \$0,%eax     jg nextdigit      subl \$0x0,%esp     pushl %edi     call printf     popl %edi     movl %ebp,%esp     ret  # compile and run under Mac OS # gcc -gstabs -o digitmac digitmac.s # ./digitmac</pre>		

digitmac.s

Fredag Mars 12, 2010